

Load the dataset from keras

```
import tensorflow as tf

# Load the dataset
mnist = tf.keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels) =
mnist.load_data()
```

Starting out

```
import numpy as np

# Initialize arrays for each digit
array_0 = []
array_1 = []
array_2 = []
array_3 = []
array_4 = []
array_5 = []
array_6 = []
array_7 = []
array_8 = []
array_9 = []

# Iterate over each image and label in the training set
for image, label in zip(train_images, train_labels):
    # Reshape the image to a column vector
    image_vector = image.reshape(784, 1)

    # Append the image to the corresponding array based on the label
    if label == 0:
        array_0.append(image_vector)
    elif label == 1:
        array_1.append(image_vector)
    elif label == 2:
        array_2.append(image_vector)
    elif label == 3:
        array_3.append(image_vector)
    elif label == 4:
        array_4.append(image_vector)
    elif label == 5:
        array_5.append(image_vector)
    elif label == 6:
        array_6.append(image_vector)
```

```

elif label == 7:
    array_7.append(image_vector)
elif label == 8:
    array_8.append(image_vector)
elif label == 9:
    array_9.append(image_vector)

# Convert the lists to NumPy arrays
array_0 = np.hstack(array_0)
array_1 = np.hstack(array_1)
array_2 = np.hstack(array_2)
array_3 = np.hstack(array_3)
array_4 = np.hstack(array_4)
array_5 = np.hstack(array_5)
array_6 = np.hstack(array_6)
array_7 = np.hstack(array_7)
array_8 = np.hstack(array_8)
array_9 = np.hstack(array_9)

# Verify the shapes of the arrays
print("Array shapes:")
print("array_0:", array_0.shape)
print("array_1:", array_1.shape)
print("array_2:", array_2.shape)
print("array_3:", array_3.shape)
print("array_4:", array_4.shape)
print("array_5:", array_5.shape)
print("array_6:", array_6.shape)
print("array_7:", array_7.shape)
print("array_8:", array_8.shape)
print("array_9:", array_9.shape)

Array shapes:
array_0: (784, 5923)
array_1: (784, 6742)
array_2: (784, 5958)
array_3: (784, 6131)
array_4: (784, 5842)
array_5: (784, 5421)
array_6: (784, 5918)
array_7: (784, 6265)
array_8: (784, 5851)
array_9: (784, 5949)

```

For each of this matrix $D(n)$ write a routine to extract the first 4 left singular vectors

(i.e., the principal components). You may call function provided by numpy. Recall that

these are orthonormal basis vectors of the column space of $D(n)$. Notice that we have used the name `array_n` in place of $D(n)$.

```
import matplotlib.pyplot as plt

# Prepare to store the singular vectors for each digit
singular_vectors = {}

# Process each digit
for digit in range(10):
    # Get all images of the current digit using boolean filtering
    digit_images = train_images[train_labels == digit]

    # Reshape images to vectors (28*28 = 784)
    digit_vectors = digit_images.reshape(digit_images.shape[0], -1).T

    # Compute the Singular Value Decomposition
    U, S, Vt = np.linalg.svd(digit_vectors, full_matrices=False)

    # Store the first 4 left singular vectors
    singular_vectors[digit] = U[:, :4]

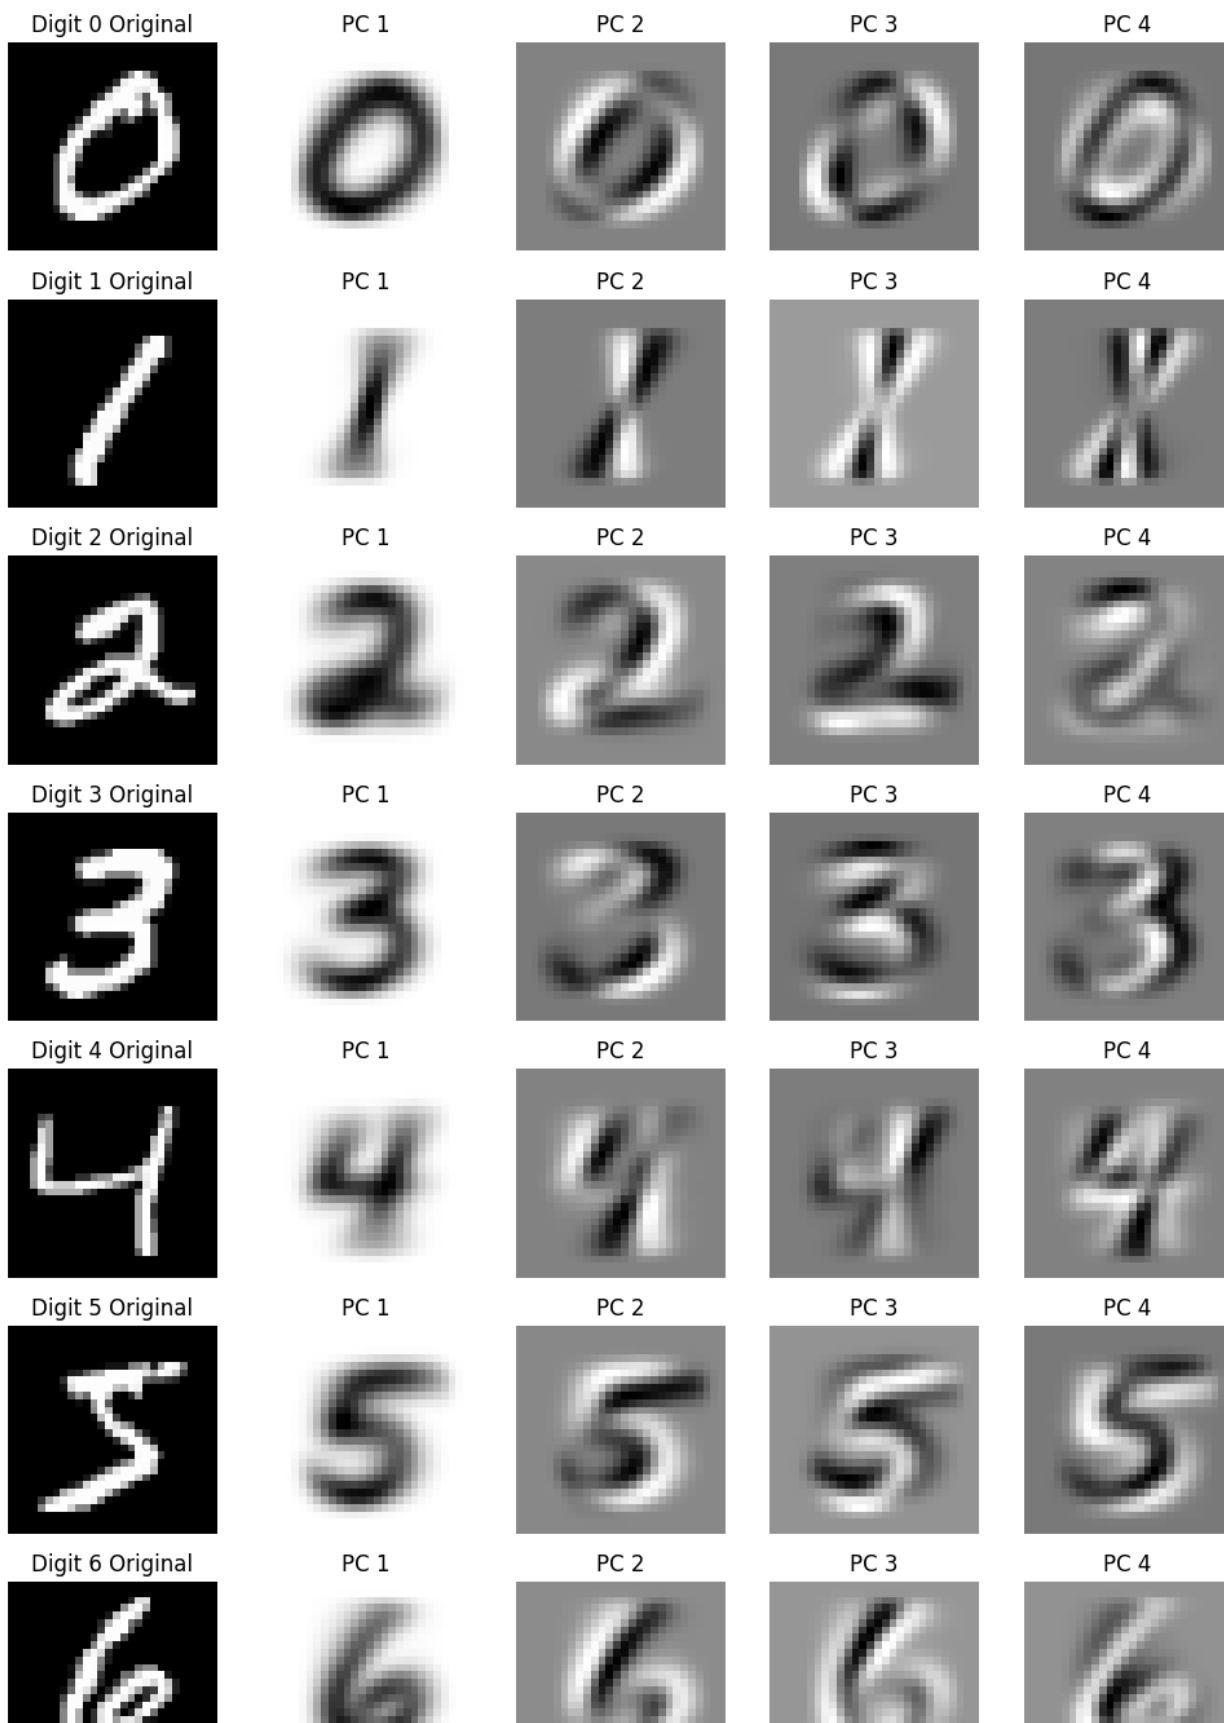
# Display the first image from the training data and the 4 principal
# component images for each digit
fig, axs = plt.subplots(10, 5, figsize=(10, 20))

for digit in range(10):
    # Display the first image from the training set
    first_image = train_images[train_labels == digit][0]
    axs[digit, 0].imshow(first_image, cmap='gray')
    axs[digit, 0].title.set_text(f'Digit {digit} Original')
    axs[digit, 0].axis('off')

    # Display the first 4 principal components reshaped as images
```

```
    for i in range(4):
        axs[digit, i+1].imshow(singular_vectors[digit][:,
i].reshape(28, 28), cmap='gray')
        axs[digit, i+1].title.set_text(f'PC {i+1}')
        axs[digit, i+1].axis('off')

plt.tight_layout()
plt.show()
```



```

test_images.shape
(10000, 28, 28)

def predict_digit(z, U_matrices):
    residuals = []
    for U in U_matrices:
        # Project z onto the space spanned by the columns of U
        x = np.dot(U.T, z)
        # Calculate the reconstruction from the low-rank approximation
        z_approx = np.dot(U, x)
        # Calculate the norm of the residual
        residual = np.linalg.norm(z_approx - z)
        residuals.append(residual)

    # The predicted digit is the one with the minimum residual
    predicted_digit = np.argmin(residuals)
    return predicted_digit

for digit in range(10):
    digit_images = train_images[train_labels == digit]

    # Reshape images to vectors (28*28 = 784)
    digit_vectors = digit_images.reshape(digit_images.shape[0], -1).T
    # Compute the Singular Value Decomposition
    U, S, Vt = np.linalg.svd(digit_vectors, full_matrices=False)

    if digit==0:
        u_0=U[:, :4]
    elif digit==1:
        u_1=U[:, :4]
    elif digit==2:
        u_2=U[:, :4]
    elif digit==3:
        u_3=U[:, :4]
    elif digit==4:
        u_4=U[:, :4]
    elif digit==5:
        u_5=U[:, :4]
    elif digit==6:
        u_6=U[:, :4]
    elif digit==7:
        u_7=U[:, :4]
    elif digit==8:
        u_8=U[:, :4]
    elif digit==9:
        u_9=U[:, :4]

```

```

# Collect all U matrices (the first 4 singular vectors for each digit)
into a list
U_matrices = [u_0, u_1, u_2, u_3, u_4, u_5, u_6, u_7, u_8, u_9]
# Reshape and predict each test image
predicted_labels = []
for test_image in test_images:
    # Flatten the test image into a vector
    test_vector = test_image.flatten()

    # Use the predict_digit function to determine the digit
    predicted_digit = predict_digit(test_vector, U_matrices)
    predicted_labels.append(predicted_digit)

# Calculate accuracy
correct_predictions = np.sum(predicted_labels == test_labels)
total_predictions = len(test_labels)
accuracy = correct_predictions / total_predictions
print(f'Accuracy of the PCA model: {accuracy:.2%}')

import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix

# Calculate the confusion matrix and per-digit accuracy
cm = confusion_matrix(test_labels, predicted_labels)
per_digit_accuracy = cm.diagonal() / cm.sum(axis=1)

# Calculate the overall average accuracy
average_accuracy = np.mean(per_digit_accuracy)

# Plotting per-digit accuracy using a line plot
plt.figure(figsize=(10, 6))
plt.plot(range(10), per_digit_accuracy, marker='o', linestyle='--',
color='b', label='Digit Accuracy')
plt.xlabel('Digits')
plt.ylabel('Accuracy')
plt.title('Per-Digit Accuracy of the PCA Model')
plt.xticks(range(10), [str(digit) for digit in range(10)])
plt.grid(True)

# Draw a horizontal line for the average accuracy
plt.axhline(y=average_accuracy, color='r', linestyle='--',
label=f'Average Accuracy ({average_accuracy:.2f})')
plt.legend()


# Show where the average line cuts the y-axis
plt.annotate(f'Average: {average_accuracy:.2%}',
xy=(0, average_accuracy),

```

```
xytext=(-0.5, average_accuracy + 0.05),
arrowprops=dict(facecolor='black', shrink=0.05),
horizontalalignment='right',
verticalalignment='center')
```

```
plt.show()
```

Accuracy of the PCA model: 91.54%



Confusion matrix

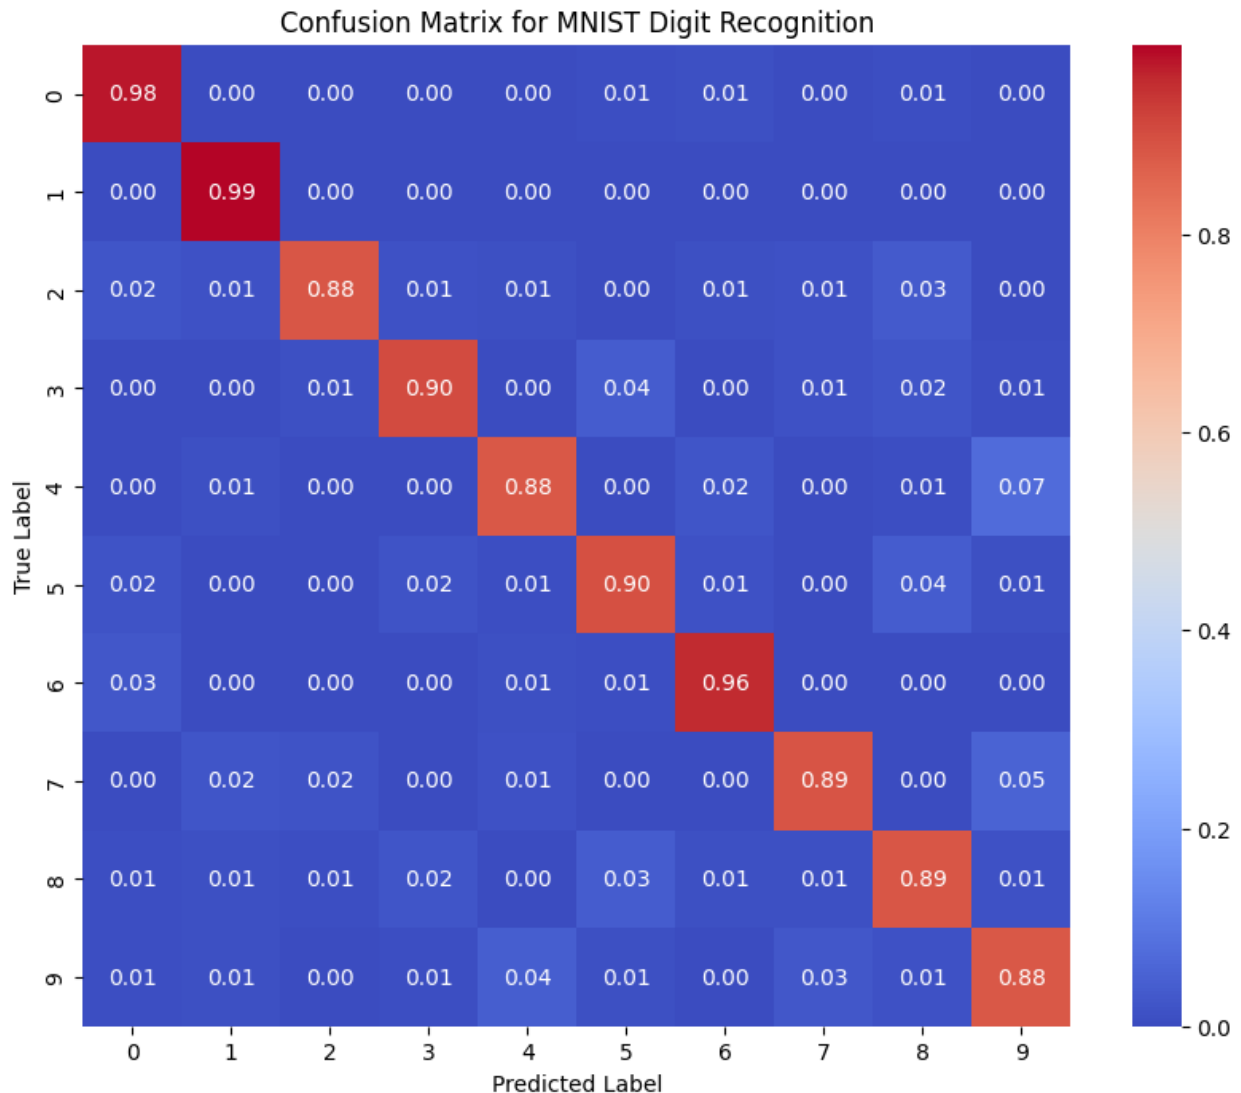
```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Calculate the confusion matrix
cm = confusion_matrix(test_labels, predicted_labels)

# Normalize the confusion matrix by the number of samples in each
class
cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
plt.figure(figsize=(10, 8))
```



```
sns.heatmap(cm_normalized, annot=True, fmt=".2f", cmap="coolwarm",
cbar=True, xticklabels=range(10), yticklabels=range(10))
plt.title('Confusion Matrix for MNIST Digit Recognition')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



Classification report Heatmap

```
from sklearn.metrics import classification_report
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
```

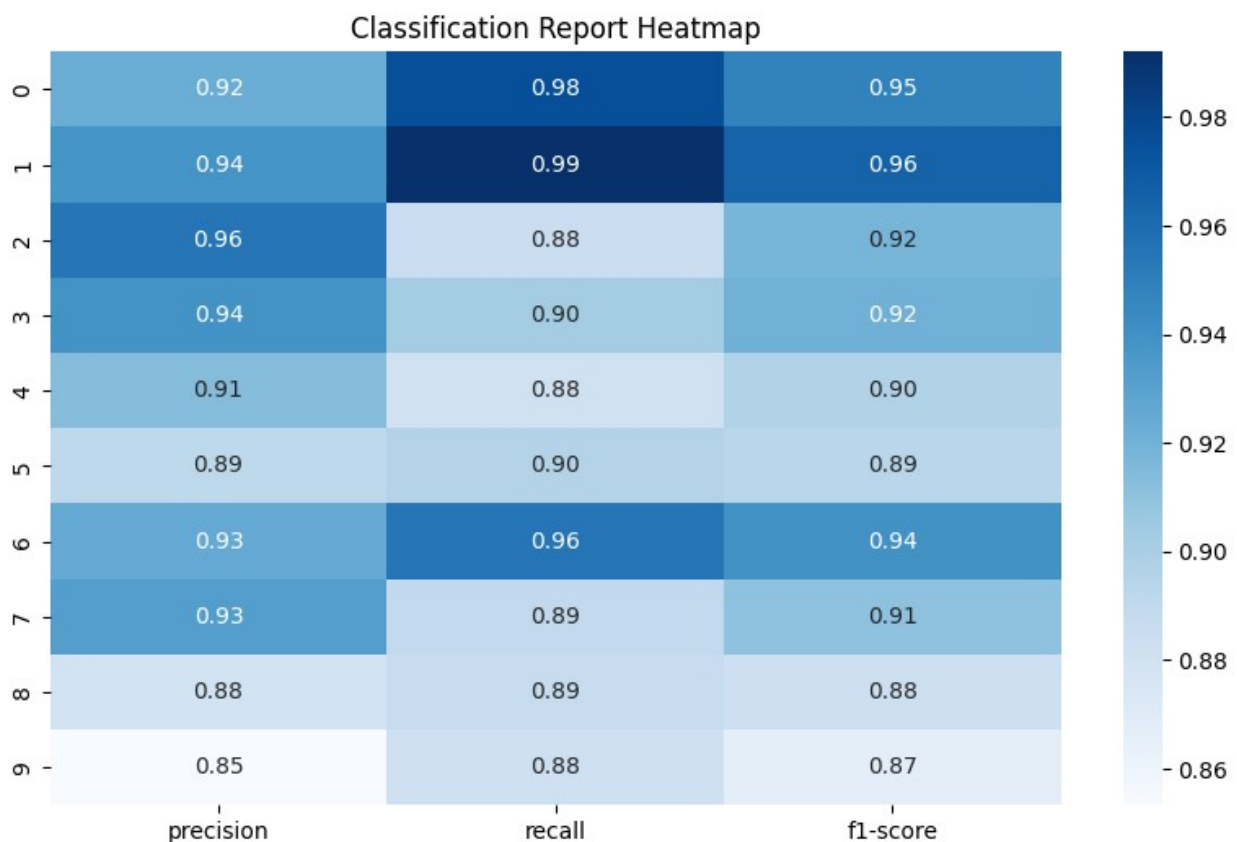
```

# Generate a classification report
report = classification_report(test_labels, predicted_labels,
                              output_dict=True)

# Convert the report to a DataFrame
df_report = pd.DataFrame(report).transpose()

# Heatmap of the classification report
plt.figure(figsize=(10, 6))
sns.heatmap(df_report.iloc[:-3, :-1], annot=True, cmap="Blues",
            fmt=".2f") # Exclude the last 3 rows and the last column
plt.title('Classification Report Heatmap')
plt.show()

```



Displaying the required 50 images.

```

# Display the first image from the training data and the 4 principal
component images for each digit
fig, axs = plt.subplots(10, 5, figsize=(10, 20))

for digit in range(10):
    # Display the first image from the training set

```

```
first_image = train_images[train_labels == digit][0]
axs[digit, 0].imshow(first_image, cmap='gray')
axs[digit, 0].title.set_text(f'Digit {digit} Original')
axs[digit, 0].axis('off')

# Display the first 4 principal components reshaped as images
for i in range(4):
    axs[digit, i+1].imshow(singular_vectors[digit][:,
i].reshape(28, 28), cmap='gray')
    axs[digit, i+1].title.set_text(f'PC {i+1}')
    axs[digit, i+1].axis('off')

plt.tight_layout()
plt.show()
```

